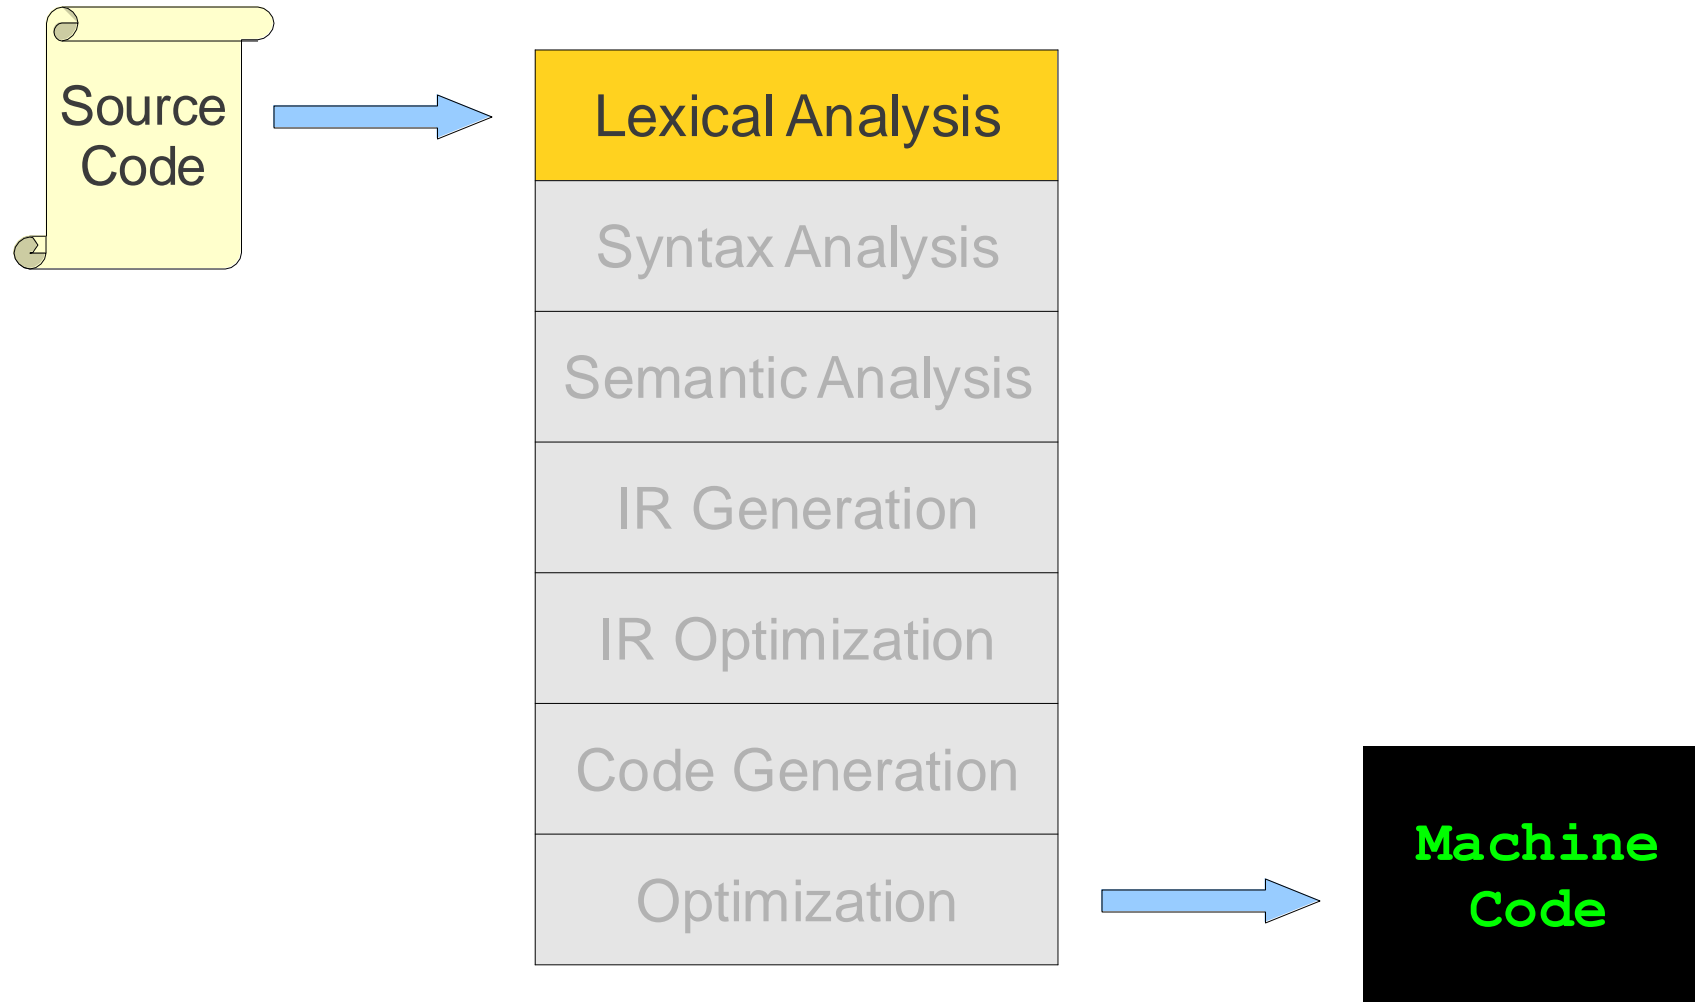


Compiler

– 2-2. Lexical Analysis –

JIEUNG KIM

Where are we?



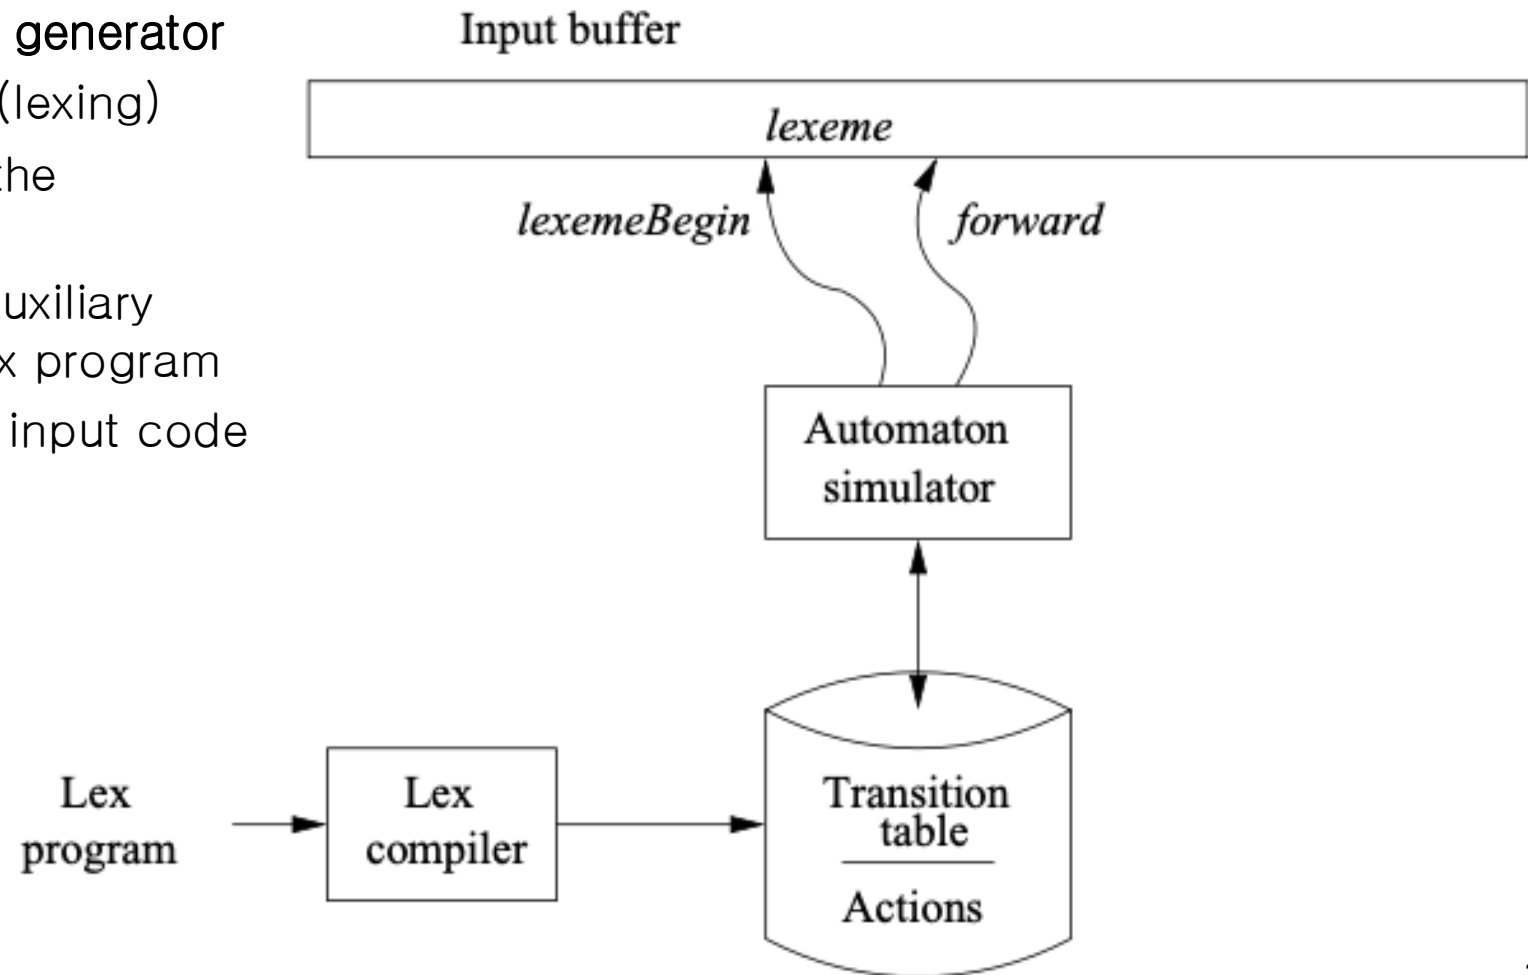
Outlines

- Basic concepts of formal grammars
- Role of the lexical analyzer
- Choose a token
- Finite automata
- Regular expression
- Specification of tokens
- Recognition of tokens
- Challenges in scanning
- Error handling
- Lex : lexical analyzer generator



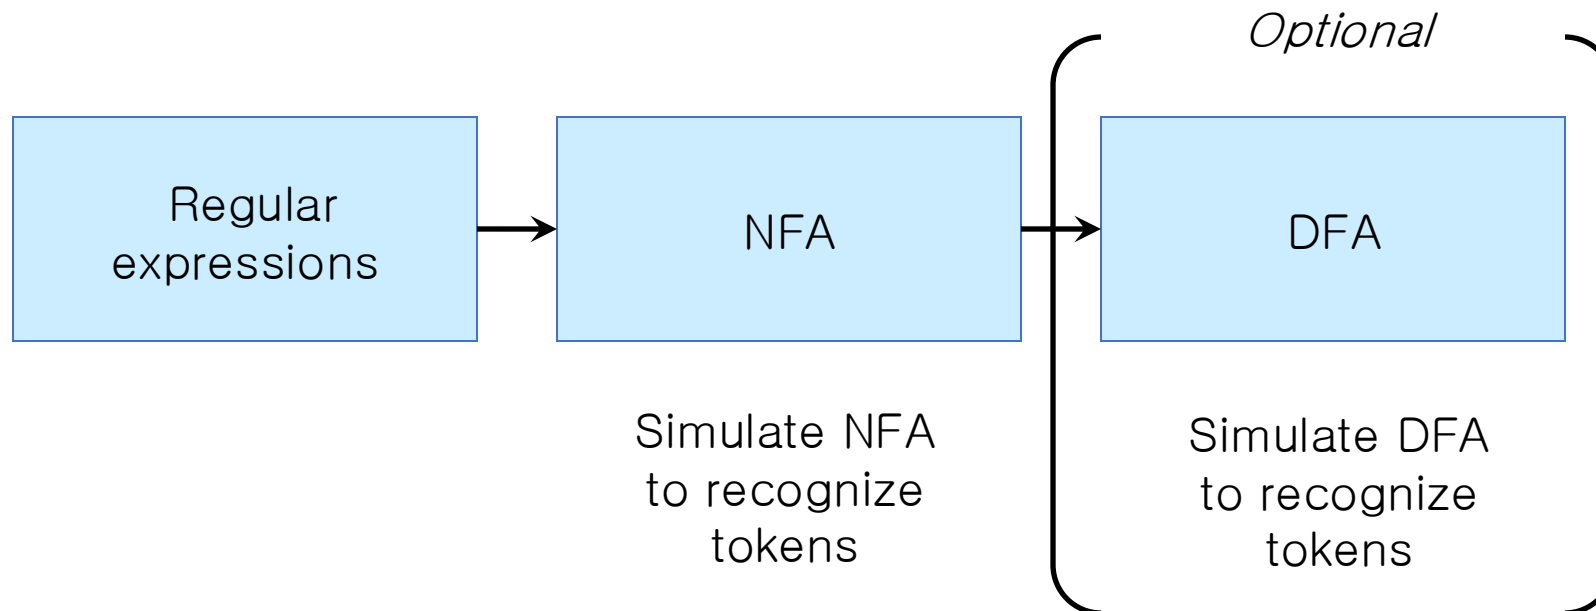
Design of lexical-analyzer generator

- Design of lexical-analyzer generator
 - Automaton simulator (lexing)
 - A transition table for the automaton
 - Directly passed auxiliary functions from lex program
 - The actions from input code



Design of a lexical analyzer generator

- Design of lexical-analyzer generator
 - Translate regular expressions to NFA
 - Translate NFA to an efficient DFA



Specification of tokens

(lexical specification of programming languages)

Specification of tokens

- Specification of tokens
 - Specifying all lexeme patterns is not efficient
 - Use Regular Expressions
- Example

```
num → digits ( . Digits ) ? ( E ( + | - ) ? digits ) ?  
digit → [ 0 - 9 ]  
digits → digit +
```


Specification of tokens

Grammars

```
stmt → if expr then stmt
      | if expr then stmt else stmt
      | ε
expr → term relop term
      | term
term → id
      | num
```

Regular definitions

```
if → if
then → then
else → else
relop → < | > | <= | >= | = | <>
id → letter ( letter | digit )*
num → digits ( . digit )* ( E ( + | - )* digits )?
letter → [A-Za-z]
digit → [0-9]
digits → digit+
```



Recognition of tokens

(using finite automata to recognize regular expressions)

Recognition of tokens

- Recognition of tokens
 - Finite automata can be used to recognize strings generated by regular expressions
 - Can build by hand or automatically
 - Not totally straightforward, but can be done systematically
 - Tools like Lex, Flex, Jlex, ANTLR do this automatically, given a set of REs for tokens



Recognition of tokens

- Recognition of tokens

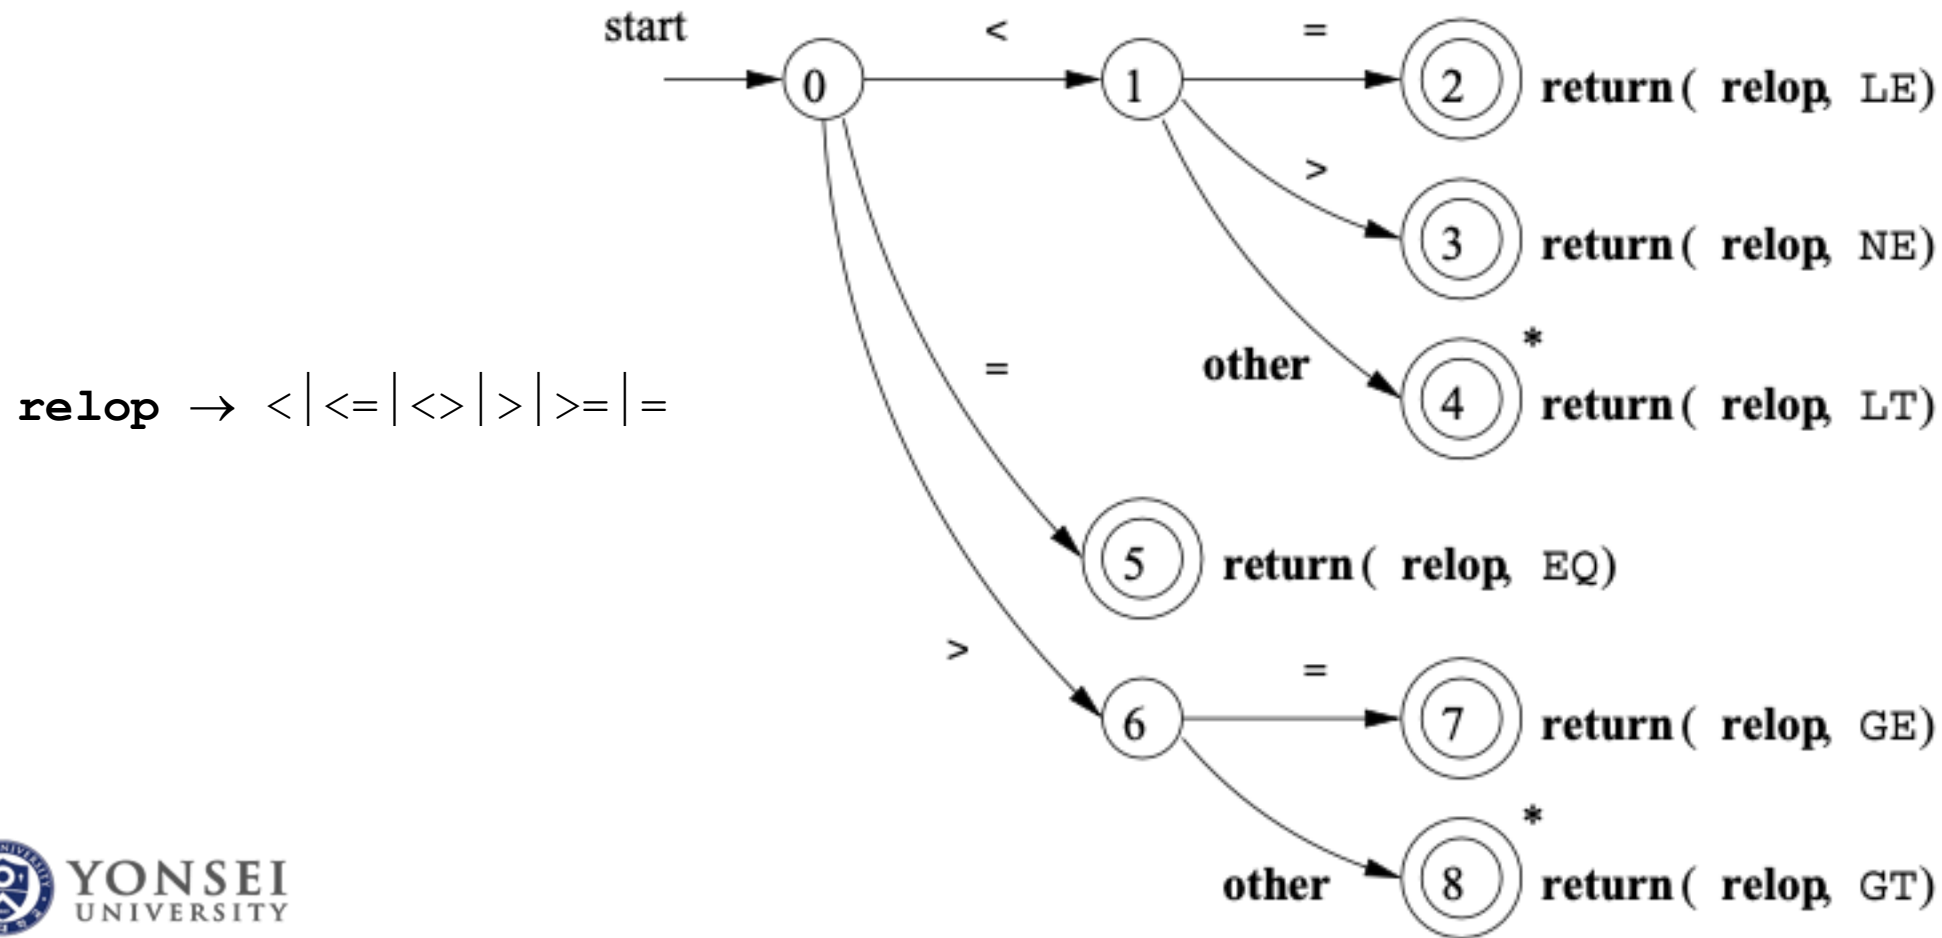
```
if → if
then → then
else → else
relop → < | > | <= | >= | = | <>
id → letter ( letter | digit ) *
num → digits ( . digits ) ? ( E ( + | - ) ? digits ) ?
letter → [A-Za-z]
digit → [0-9]
digits → digit +
```

The issues are “*how to recognize the keywords– if, then else, relop, id, number*”




Recognition of tokens

- Transition diagram for relop



Recognition of tokens

- Transition diagram for unsigned numbers





Recognition of tokens

- Transition diagram for reserved keywords and IDs
 - Two ways to handle reserved word
 - Install keyword symbol table initially
 - Hand-written scanner: look up identifier-like things in table of keywords to classify (good application of perfect hashing)
 - Create separate diagram for each keyword
 - Lots of states, but efficient (no extra lookup step)
 - Machine-generated scanner: generate DFA with appropriate transitions to recognize keywords




Recognition of tokens

- Transition diagram for reserved words and IDs



Recognition of tokens

- Transition diagram for white spaces



Simple hand-written scanner example

- Example: DFA and hand-written scanner (ad hoc scanners)
 - Idea: show a hand-written DFA for some typical programming language constructs
 - Then use to construct hand-written scanner
 - Setting: scanner is called whenever the parser needs a new token
 - Scanner stores current position in input
 - Starting there, use a DFA to recognize the longest possible input sequence that makes up a token and return that token



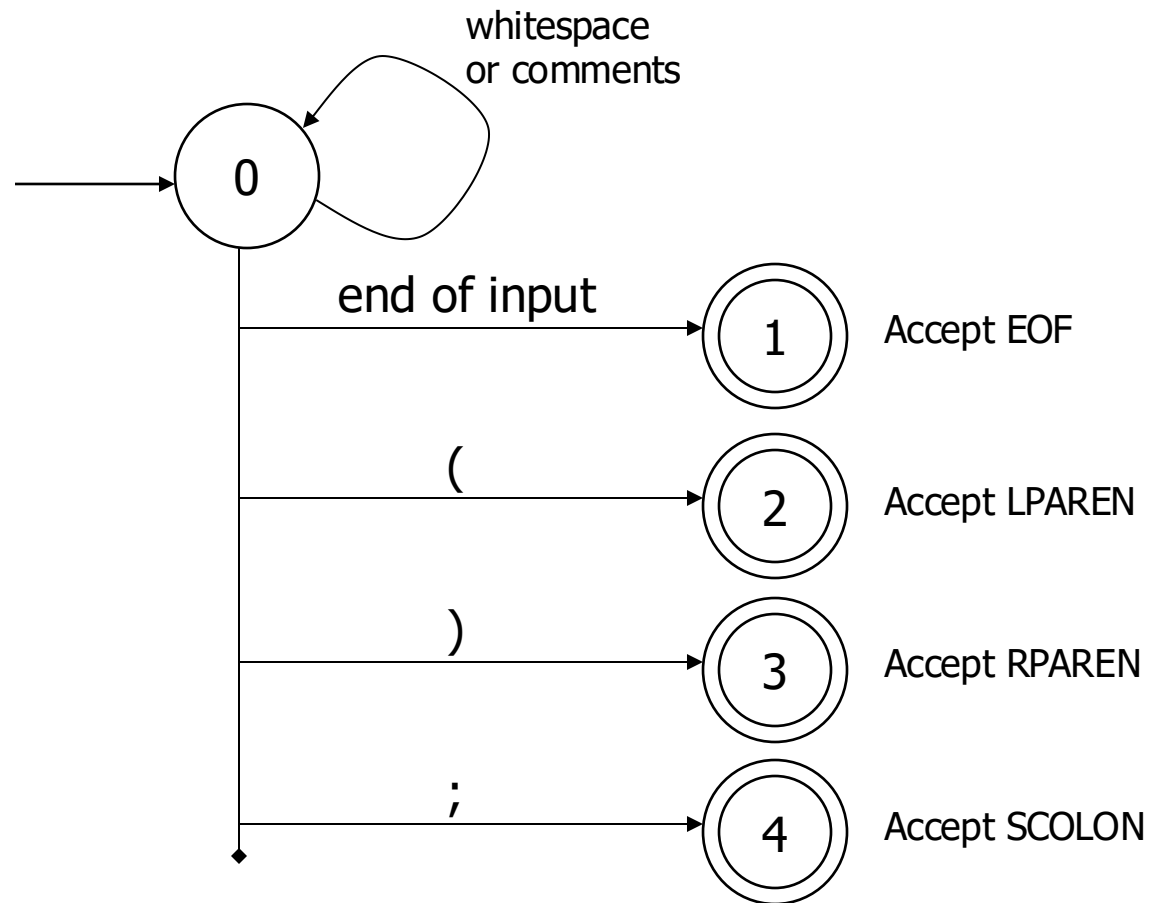
Simple hand-written scanner example

- Example: DFA for hand-written scanner (ad hoc scanners)
 - Loop and switch scanner
 - Big nested switch/case statements
 - Lots of `gets()` / `ungetc()` calls
 - Must be error-prone, use only if
 - The lexical structure of the language is very simple
 - The tools do not provide what you need for your token definition
 - Changing keyword is problematic
 - Very difficult to show correctness



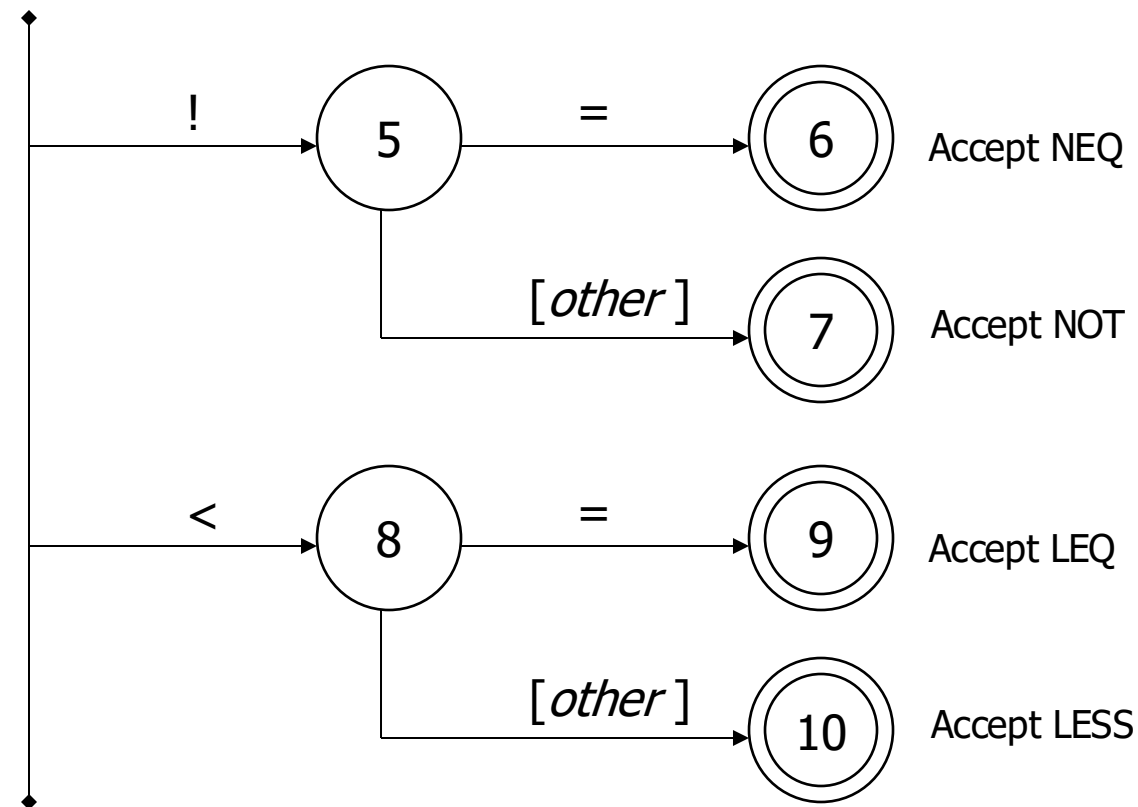
Simple hand-written scanner example

- Example 2 (1)



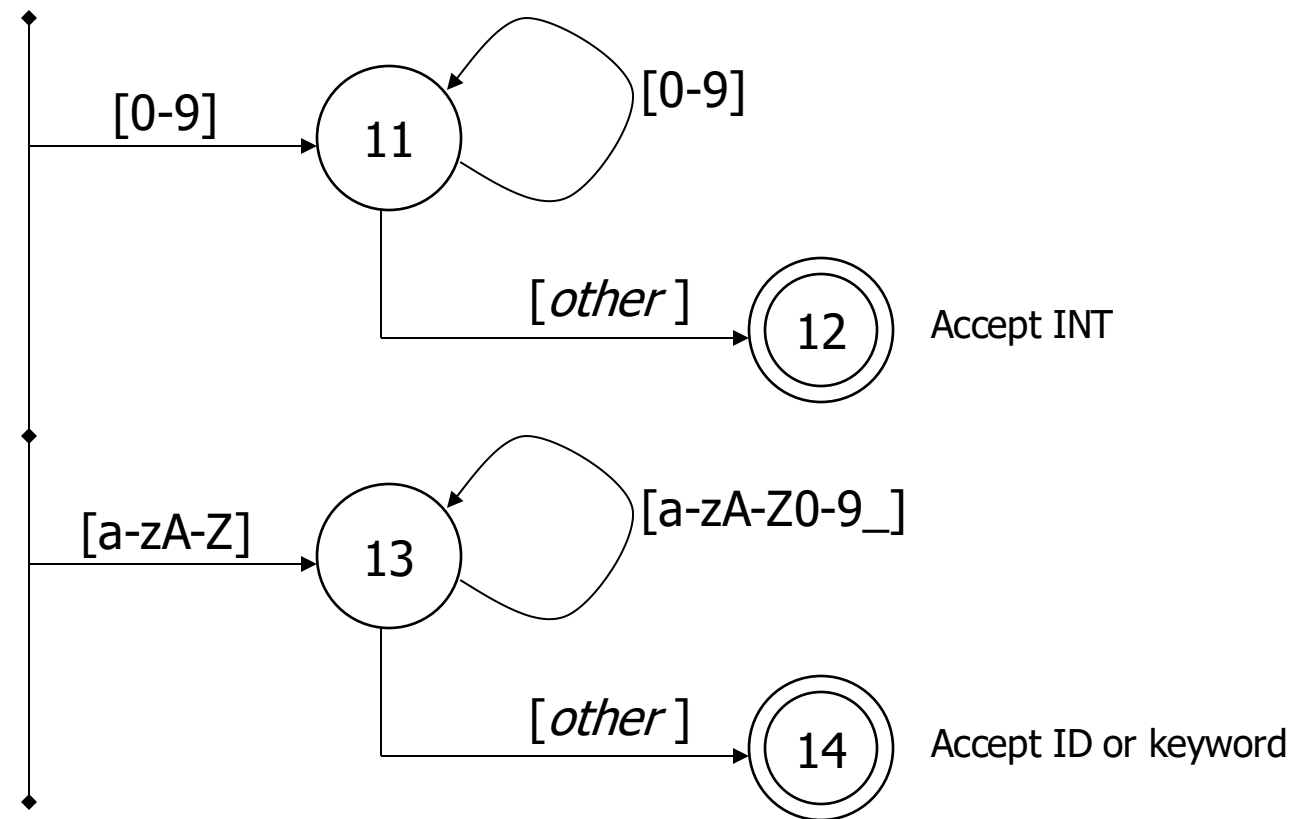
Simple hand-written scanner example

- Example 2 (3)



Simple hand-written scanner example

- Example 2 (4)



Simple hand-written scanner example

- Implementing a scanner by hand – token representation

```
public class Token {  
    public int kind;           // token's lexical class  
    public int intVal;        // integer value if class = INT  
    public String id;         // actual identifier if class = ID  
    // lexical classes  
    public static final int EOF = 0;    // "end of file" token  
    public static final int ID = 1;     // identifier, not keyword  
    public static final int INT = 2;    // integer  
    public static final int LPAREN = 4;  
    public static final int SCOLN = 5;  
    public static final int WHILE = 6;  
    // etc. etc. etc. ...  
}
```



Simple hand-written scanner example

- Implementing a scanner by hand – read inputs

```
// global state and methods

static char nextch;    // next unprocessed input character

// advance to next input char
void getch() { ... }

// skip whitespace and comments
void skipWhitespace() { ... }
```



Simple hand-written scanner example

- Implementing a scanner by hand – scanner getToken() method

```
// return next input token
public Token getToken() {
    Token result;

    skipWhiteSpace();

    if (no more input) {
        result = new Token(Token.EOF);
        return result;
    }

    switch(nextch) {
        case '(': result = new Token(Token.LPAREN); getch(); return result;
        case ')': result = new Token(Token.RPAREN); getch(); return result;
        case ';': result = new Token(Token.SCOLON); getch(); return result;

        // etc. ...
    }
}
```



Simple hand-written scanner example

- Implementing a scanner by hand – scanner getToken() method

```
case '!':                                // ! or !=
    getch();
    if (nextch == '=') {
        result = new Token(Token.NEQ); getch(); return result;
    } else {
        result = new Token(Token.NOT); return result;
    }

case '<':                                // < or <=
    getch();
    if (nextch == '=') {
        result = new Token(Token.LEQ); getch(); return result;
    } else {
        result = new Token(Token.LESS); return result;
    }
// etc. ...
```



Simple hand-written scanner example

- Implementing a scanner by hand – scanner getToken() method

```
case '0': case '1': case '2': case '3': case '4':  
case '5': case '6': case '7': case '8': case '9':  
    // integer constant  
    String num = nextch;  
    getch();  
    while (nextch is a digit) {  
        num = num + nextch; getch();  
    }  
    result = new Token(Token.INT, Integer(num).intValue());  
    return result;  
...
```



Simple hand-written scanner example

- Implementing a scanner by hand – scanner getToken() method

```
case 'a': ... case 'z':  
case 'A': ... case 'Z':  
    // id or keyword  
    string s = nextch; getch();  
    while (nextch is a letter, digit, or underscore) {  
        s = s + nextch; getch();  
    }  
    if (s is a keyword) {  
        result = new Token(keywordTable.getKind(s));  
    } else {  
        result = new Token(Token.ID, s);  
    }  
    return result;
```



Challenges in scanning

Challenges in scanning

- Challenges
 - How do we determine which lexemes are associated with each token?
 - When there are multiple ways we could scan the input, how do we know which one to pick?
 - How do we address these concerns efficiently?



Challenges in scanning

- Lexing ambiguity

T_For

T_ID

for

[A-Za-z_] [A-Za-z0-9_]*

f	o	r	t
---	---	---	---

f	o	r	t
f	o	r	t
f	o	r	t
f	o	r	t
f	o	r	t

f	o	r	t
f	o	r	t
f	o	r	t
f	o	r	t



Challenges in scanning

- Conflict resolution
 - How do we determine which lexemes are associated with each token?
 - When there are multiple ways we could scan the input, how do we know which one to pick?
 - How do we address these concerns efficiently?
 - Assume all tokens are specified as regular expressions.
 - Algorithm: **Left-to-right scan**.
 - Tiebreaking rule one: **Maximal munch**.
 - Always match the longest possible prefix of the remaining text.



Challenges in scanning

- Lexing ambiguity

T_For

T_ID

for

[A-Za-z_] [A-Za-z0-9_]*

f	o	r	t
---	---	---	---

f	o	r	t
---	---	---	---




Challenges in scanning

- Implementing maximal munch
 - Given a set of regular expressions, how can we use them to implement maximum munch?
 - Idea:
 - Convert expressions to NFAs.
 - Run all NFAs in parallel, keeping track of the last match.
 - When all automata get stuck, report the last match and restart the search at that point.




Challenges in scanning

- Implementing maximal munch




Challenges in scanning

- Implementing maximal munch




Challenges in scanning

- Implementing maximal munch




Challenges in scanning

- Implementing maximal munch




Challenges in scanning

- Implementing maximal munch



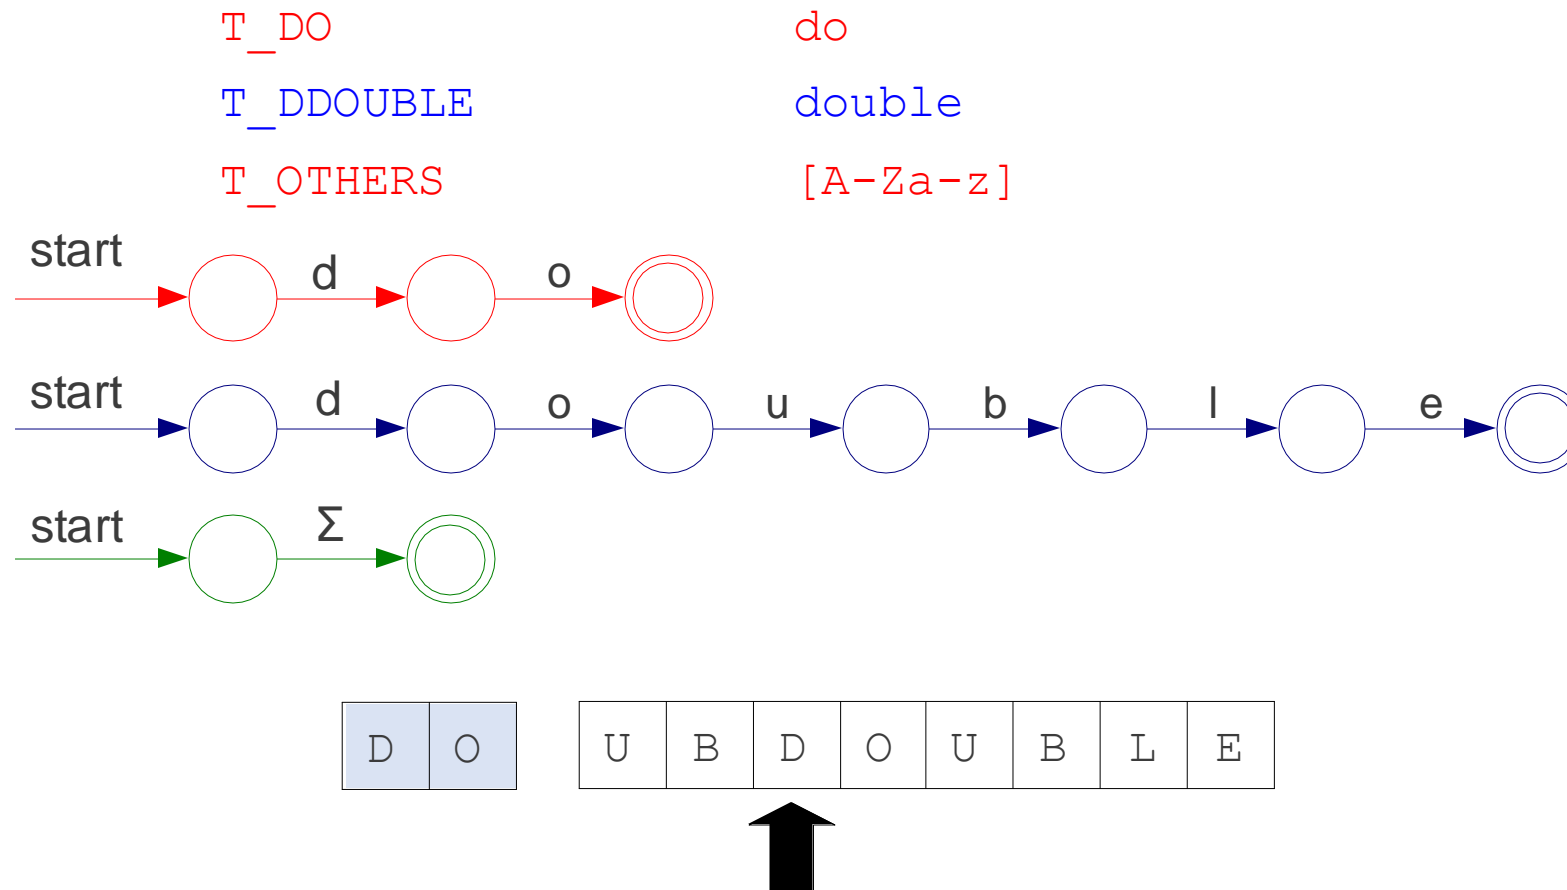
Challenges in scanning

- Implementing maximal munch



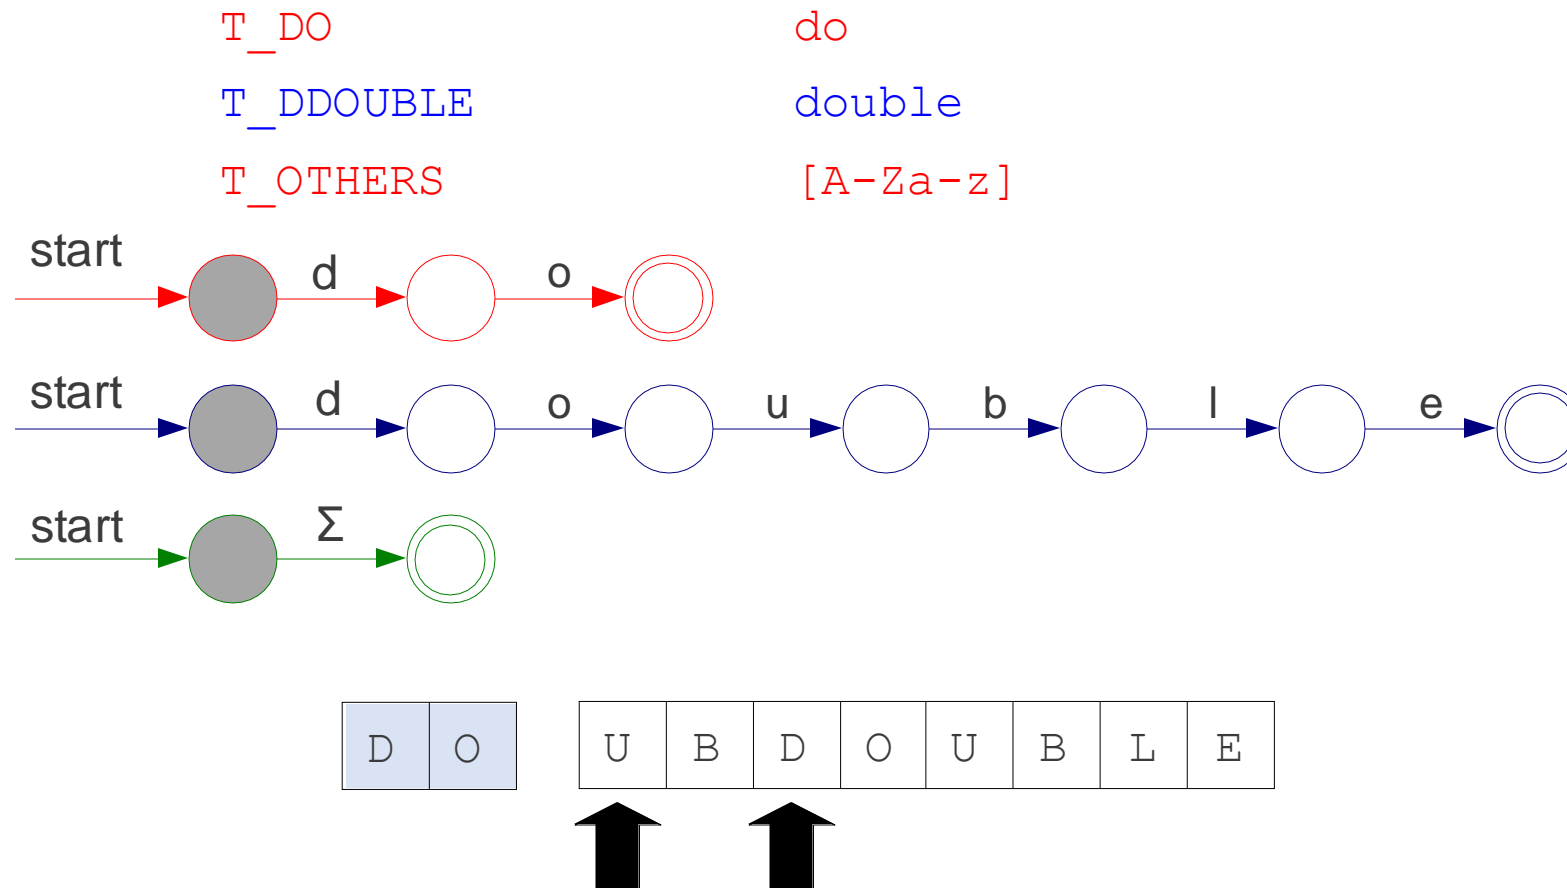
Challenges in scanning

- Implementing maximal munch



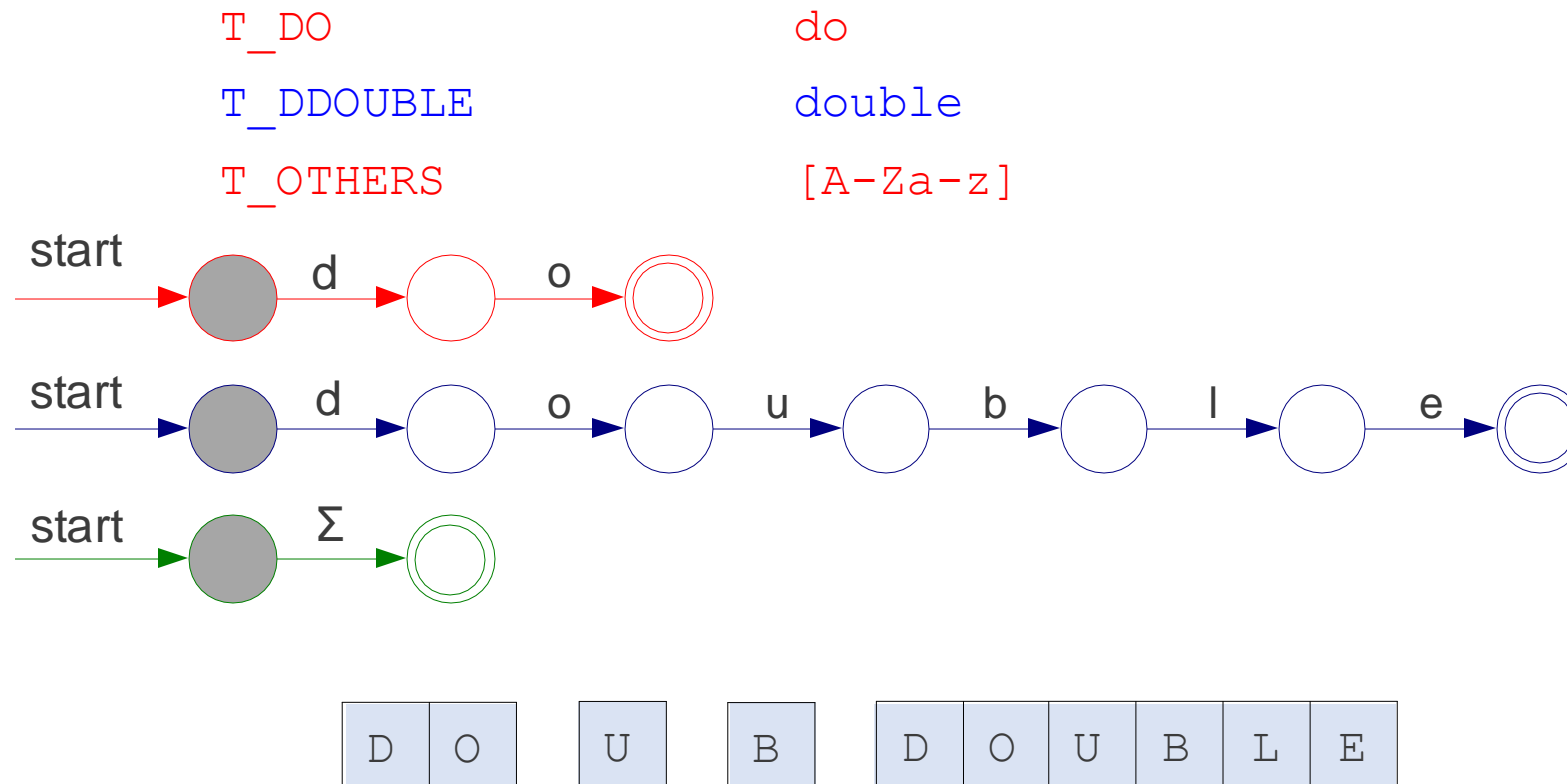
Challenges in scanning

- Implementing maximal munch



Challenges in scanning

- Implementing maximal munch



Challenges in scanning

- Other conflicts
 - When two regular expressions apply, choose the one with the greater “priority.”
 - Simple priority system: **Pick the rule that was defined first.**

T_DO

do

T_DDOUBLE

double

T_ID

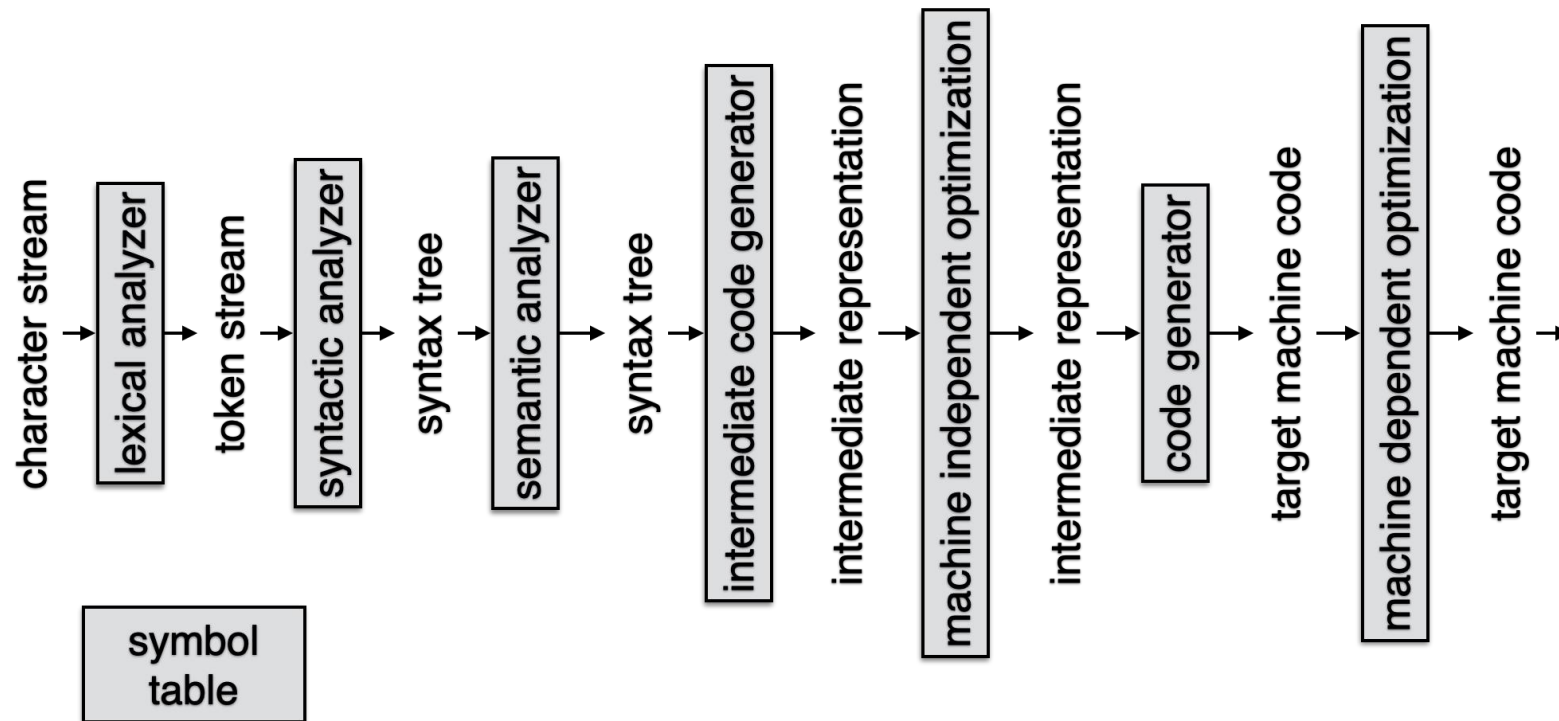
[A-Za-z_] [A-Za-z0-9_]*

D	O	U	B	L	E
---	---	---	---	---	---

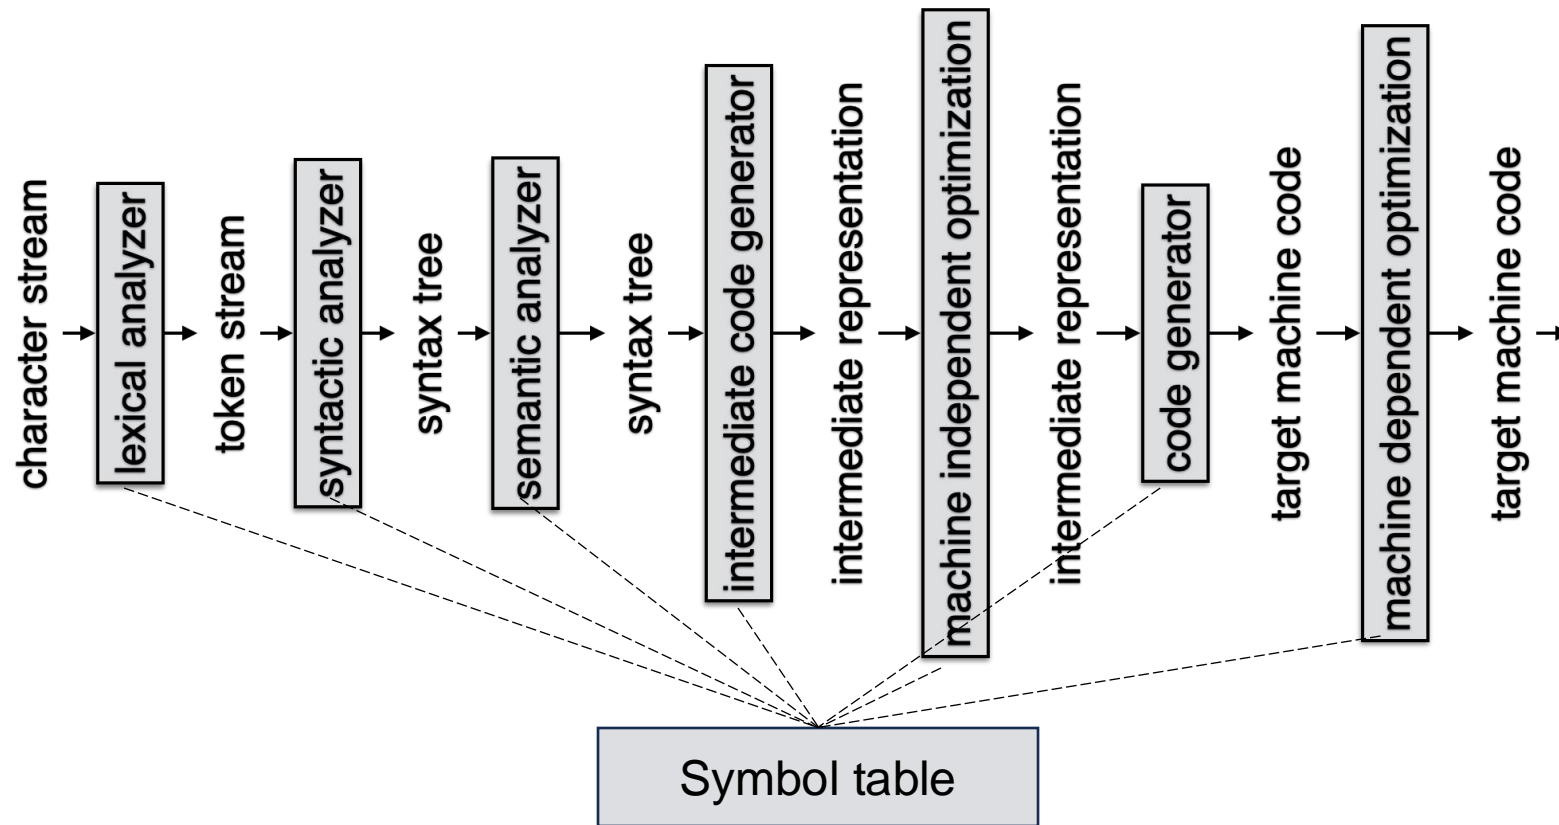


Error handling

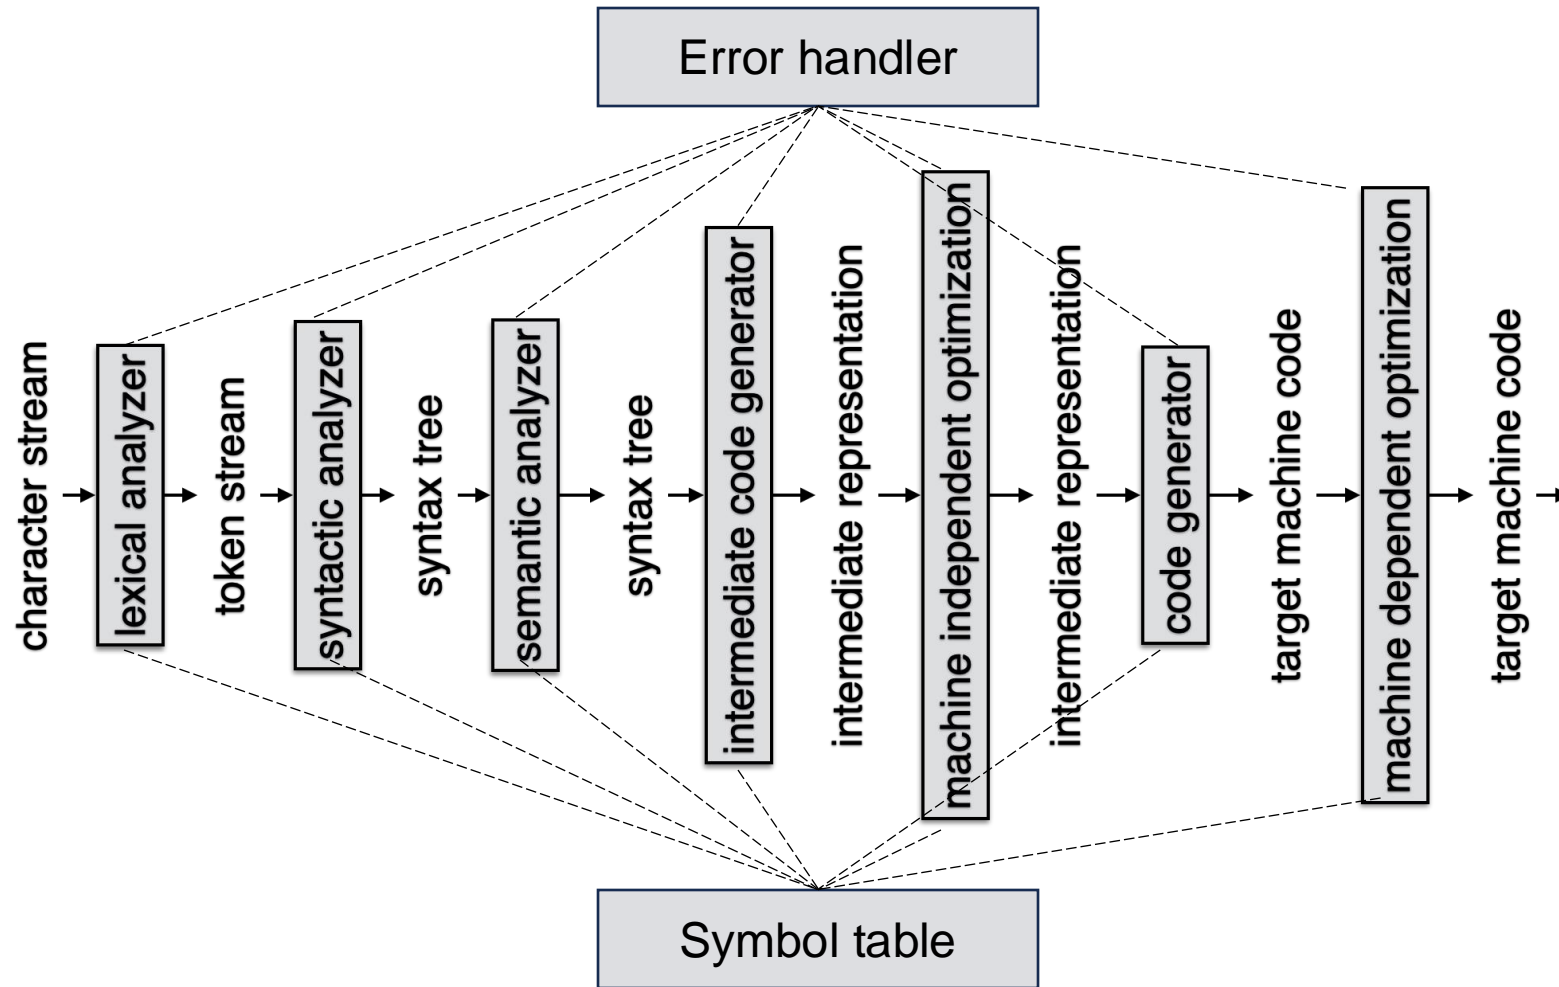
Compiler architecture



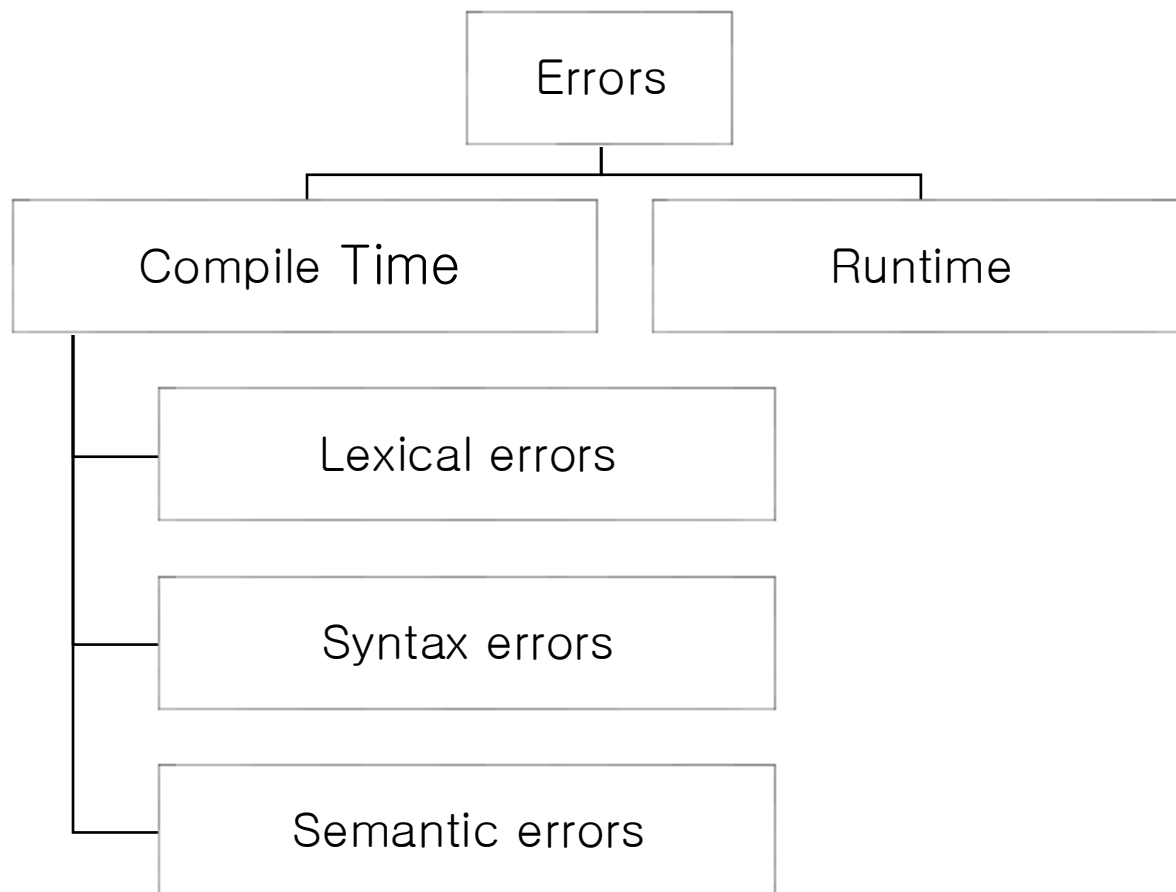
Compiler architecture



Compiler architecture

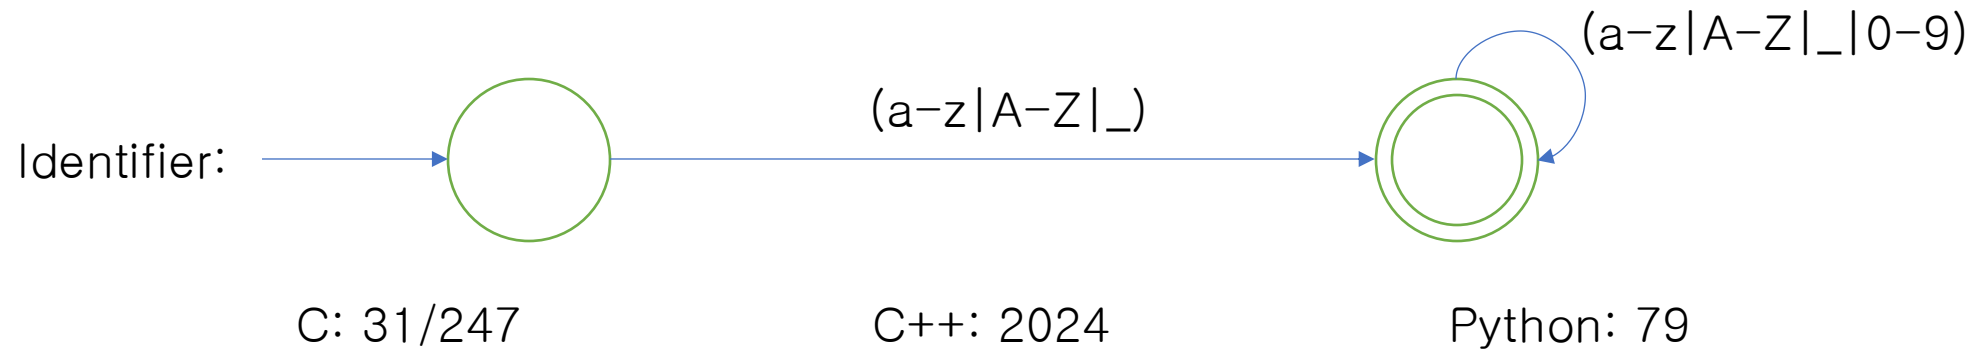


Kinds of errors



Lexical errors

- Identifiers that are way too long



- Exceeding length of numeric constants

unsinged int i = 4294967300;

Size: 4 bytes 0 ~ 4,294,967,295



Lexical errors

- Numeric constants which are ill-formed

```
int i = 4567$91;
```

- Illegal characters that are absent from the source code.

```
char x[] = "YONSEI UNIVERSITY";$
```

Lexical error recovery

- Panic-mode recovery
 - Successive characters from the input are removed one at a time until a designated set of synchronizing tokens is found.
 - Synchronizing tokens are delimiters such as ; or }
 - The advantage is that it is easy to implement and guarantees not to go into an infinite loop
 - The disadvantage is that a considerable amount of input is skipped without checking it for additional errors

int i = 4YONSEI;

```
while (condition)
{
  -----
  -----
  -----
}
```



Lexical error recovery

- Transpose of two adjacent characters.

```
unoin test  
{  
  int x;  
  float y;  
} T1;
```



```
union test  
{  
  int x;  
  float y;  
} T1;
```

- Insert a missing character.

```
it YONSEI;
```



```
int YONSEI;
```

- Delete an unknown or extra character.

```
intt YONSEI;
```



```
int YONSEI;
```

- Replace a character with another.

```
itt YONSEI;
```



```
int YONSEI;
```



Lex : lexical analyzer generator




Lex : lexical analyzer generator

- The Lex and Flex scanner generators
 - *Lex* and its newer cousin *flex* are scanner generators
 - Systematically translate regular definitions into C source code for efficient scanning
 - Generated code is easy to integrate in C applications



Lex : lexical analyzer generator

- Create a lexical analyzer with Lex and Flex



Lex programs

- Form
 - Declarations : declare variables, manifest constants and regular definitions
 - Translation rules : Pattern {Action}
 - Pattern is regular expression
 - Auxiliary functions : whatever additional functions

```
Declarations
%%
Translation rules
%%
Auxiliary functions
```


Lex programs

- Co-work with parser
 - Parser calls lexical analyzer
 - It finds longest prefix of input that matches the pattern P_i and it executes action A_i
 - A_i returns to the parser.
 - If it does not (P_i is whitespace or comments), lexical analyzer proceed to next lexemes
 - Lexical Analyzer returns one value, **token name** to parser.



Lex programs

- Example 1

Translation rules

```
%{
#include <stdio.h>
}%
%%
[0-9]+  { printf("%s\n", yytext); }
.|\\n   { }
%%
main()
{
    yylex();
}
```

Contains the matching lexeme

Invokes the lexical analyzer

```
lex spec.l
gcc lex.yy.c -ll
./a.out < spec.l
```



Lex programs

- Example 2

Translation
rules

```
%{
#include <stdio.h>
int ch = 0, wd = 0, nl = 0;
}%
delim      [ \t]+
%%
\n         { ch++; wd++; nl++; }
^{delim}   { ch+=yyleng; }
{delim}    { ch+=yyleng; wd++; }
.          { ch++; }
%%
main() {
    yylex();
    printf("%8d%8d%8d\n", nl, wd, ch);
}
```

Regular
definition



Lex programs

- Example 3

Translation rules

```
%{
#include <stdio.h>
%}
digit      [0-9]
letter     [A-Za-z]
id         {letter}({letter}|{digit})*
%%
{digit}+   { printf("number: %s\n", yytext); }
{id}       { printf("ident: %s\n", yytext); }
.          { printf("other: %s\n", yytext); }
%%
main() {
    yylex();
}
```

Regular definitions



Lex programs

- Conflict resolution and lookahead operator in Lex
 - Conflict resolution
 - Input stream “if” or “<=“ ?
 - “<“ + “=“ and “<=“
 - Prefer longer prefix to the shorter one
 - Regular expressions for *if* and {id} match
 - Prefer the pattern listed first in lex program
 - The lookahead operator
 - IF(I,J) = 3 ?
 - 2-d array, named if, and assignment statement.
 - IF(condition) Then
 - Lex rule need to know if is a keyword if
 - IF / ~~W~~(.* ~~W~~) {letter}



$x = 3 + 42 * (s - t)$

A tokenizer splits the string into individual tokens

'x', '=', '3', '+', '42', '*', '(', 's', '-', 't', ')'

Tokens are usually given names to indicate what they are. For example:

'ID', 'EQUALS', 'NUMBER', 'PLUS', 'NUMBER', 'TIMES',
'LPAREN', 'ID', 'MINUS', 'ID', 'RPAREN'

More specifically, the input is broken into pairs of token types and values. For example:

('ID', 'x'), ('EQUALS', '='), ('NUMBER', '3'),
('PLUS', '+'), ('NUMBER', '42'), ('TIMES', '*'),
('LPAREN', '('), ('ID', 's'), ('MINUS', '-'),
('ID', 't'), ('RPAREN', ')')

```

import ply.lex as lex

# List of token names.  This is always required
tokens = (
    'NUMBER',
    'PLUS',
    'MINUS',
    'TIMES',
    'DIVIDE',
    'LPAREN',
    'RPAREN',
)

# Regular expression rules for simple tokens
t_PLUS = r'\+'
t_MINUS = r'\-'
t_TIMES = r'\*'
t_DIVIDE = r'\/'
t_LPAREN = r'\('
t_RPAREN = r'\)'

# A regular expression rule with some action code
def t_NUMBER(t):
    r'\d+'
    t.value = int(t.value)
    return t

# Define a rule so we can track line numbers
def t_newline(t):
    r'\n+'
    t.lexer.lineno += len(t.value)

# A string containing ignored characters (spaces and tabs)
t_ignore = ' \t'

# Error handling rule
def t_error(t):
    print("Illegal character '%s'" % t.value[0])
    t.lexer.skip(1)

# Build the lexer
lexer = lex.lex()

```



```

# Test it out
data = '''
3 + 4 * 10
  + -20 *2
'''

# Give the lexer some input
lexer.input(data)

# Tokenize
while True:
    tok = lexer.token()
    if not tok:
        break      # No more input
    print(tok)

```

When executed, the example will produce the following output:

```

$ python example.py
LexToken(NUMBER,3,2,1)
LexToken(PLUS,'+',2,3)
LexToken(NUMBER,4,2,5)
LexToken(TIMES,'*',2,7)
LexToken(NUMBER,10,2,10)
LexToken(PLUS,'+',3,14)
LexToken(MINUS,'-',3,16)
LexToken(NUMBER,20,3,18)
LexToken(TIMES,'*',3,20)
LexToken(NUMBER,2,3,21)

```



Questions?

